

Phase 2 Detailed Report: Implementation, Optimization, and Evaluation

Project: Facial Expression Recognition (FER) System

Course: Artificial Intelligence - K. N. Toosi University of Technology

Instructor: Dr. Pishgoo

Team Members: Soroush Soleimani, Parham Kootzari, Amirhossein Babaee

1. Evolutionary Architectural Design

In Phase 2, we moved beyond the baseline (Phase 1) to address the "Underfitting" problem. The complexity of facial muscle movements requires a high-capacity model to learn non-linear spatial features.

- **Deep CNN Architecture:** Our final model is a deep Convolutional Neural Network with approximately **5.8 Million parameters**. This scale was necessary to map the high-dimensional input of 48x48 pixels into 7 distinct emotion manifolds.

- **Feature Extraction Layers:** We utilized sequential convolutional blocks. Each block increases in depth (64, 128, 256, 512 filters) to capture everything from simple edges to complex facial structures like the curvature of a smile or the furrow of a brow.

- **Batch Normalization & Dropout:** To stabilize the internal covariate shift and force the network to learn redundant representations (preventing overfitting), we applied **BatchNormalization** and a high **Dropout** rate of 0.5 in the fully connected layers.

عکس

2. Advanced Training Strategy: The Turning Point

The most significant achievement in Phase 2 was the strategic control of the training process. We didn't just "run" the model; we managed its convergence.

ReduceLROnPlateau (The Key to 90% Accuracy): We observed that the model often got stuck in local minima. By implementing this callback, the system automatically reduced the Learning Rate by a factor of 0.2 whenever the validation loss stopped improving. This "fine-tuning" phase allowed the model to achieve a precision that a static learning rate could never reach.

Optimization Function: We used **Adam Optimizer**, which combines the benefits of AdaGrad and RMSProp, providing an adaptive gradient descent that is well-suited for image data.

Early Stopping: To ensure our model remains "Generalizable" and not just "Memorized," we used EarlyStopping to halt training exactly when the validation loss reached its global minimum.

3. Performance Analysis & Visualization

The training results exceeded our initial expectations, reaching a peak accuracy of ~90%.

Convergence Analysis: As shown in our training plots, the Training and Validation curves follow each other closely. This is a "Success Indicator" in Deep Learning, showing that our Data Augmentation from Phase 1 and our Regularization (Dropout) in Phase 2 worked in perfect harmony.

Loss Minimization: The categorical cross-entropy loss dropped significantly, proving that the model's "Confidence" in its predictions increased exponentially over 50+ epochs.

عکس

4. Final Evaluation & Metric Deep-Dive

Using the [evaluation.ipynb](#) notebook, we conducted a rigorous test on unseen data.

4.1. Confusion Matrix & Intraclass Challenges

The matrix shows that "Happy" and "Surprise" are almost perfectly classified. However, the model provides a fascinating insight into human psychology: it occasionally confuses "Sad" with "Neutral." This is a known challenge in the FER-2013 dataset where many "Neutral" faces have a slight downward lip curvature, mimicking "Sadness."

عکس

4.2. ROC Analysis (The Gold Standard)

Our ROC Curves are the strongest evidence of the model's power. Almost all classes achieved an **AUC (Area Under Curve) of 0.98 or higher**.

What this means: An AUC of 0.98 means that if we pick a random "Happy" image and a random "Angry" image, there is a 98% chance the model will rank them correctly. This confirms that our model's probability distribution is nearly ideal.

عکس

5. Final Conclusions

The project successfully met all industrial and academic criteria:

Resolution of Underfitting: Successfully scaled from a simple baseline to a 5.8M

parameter powerhouse.

Strategic Optimization: Proved that callbacks like ReduceLROnPlateau are essential for high-performance AI.

Deployment Readiness: With 90% accuracy and stable loss, the model is ready to be integrated into real-time applications (Webcam Demo).